

✓ Project Overview

This analysis aims to provide actionable insights to XYZ for their cab industry investment decision by evaluating the performance of two cab companies using EDA.

```
import pandas as pd
```

```
# Upload and load datasets
```

```
cab_data = pd.read_csv('/content/Cab_Data.csv')
city_data = pd.read_csv('/content/City.csv')
customer_data = pd.read_csv('/content/Customer_ID.csv')
transaction_data = pd.read_csv('/content/Transaction_ID.csv')
```

```
# Display the first few rows of each dataset
```

```
print("Cab Data:")
print(cab_data.head())
```

```
print("\nCity Data:")
print(city_data.head())
```

```
print("\nCustomer Data:")
print(customer_data.head())
```

```
print("\nTransaction Data:")
print(transaction_data.head())
```

```
↗ Cab Data:
```

	Transaction ID	Date of Travel	Company	City	KM Travelled \
0	10000011	42377	Pink Cab	ATLANTA GA	30.45
1	10000012	42375	Pink Cab	ATLANTA GA	28.62
2	10000013	42371	Pink Cab	ATLANTA GA	9.04
3	10000014	42376	Pink Cab	ATLANTA GA	33.17
4	10000015	42372	Pink Cab	ATLANTA GA	8.73

	Price Charged	Cost of Trip
0	370.95	313.635
1	358.52	334.854
2	125.20	97.632
3	377.40	351.602
4	114.62	97.776

```
City Data:
```

	City	Population	Users
0	NEW YORK NY	8,405,837	302,149
1	CHICAGO IL	1,955,130	164,468
2	LOS ANGELES CA	1,595,037	144,132
3	MIAMI FL	1,339,155	17,675
4	SILICON VALLEY	1,177,609	27,247

```
Customer Data:
```

	Customer ID	Gender	Age	Income (USD/Month)
0	29290	Male	28	10813
1	27703	Male	27	9237
2	28712	Male	53	11242
3	28020	Male	23	23327
4	27182	Male	33	8536

```
Transaction Data:
```

	Transaction ID	Customer ID	Payment_Mode
0	10000011	29290	Card
1	10000012	27703	Card
2	10000013	28712	Cash
3	10000014	28020	Cash
4	10000015	27182	Card

```
# Check for missing values
```

```
print("Missing Values in Cab Data:\n", cab_data.isnull().sum())
print("Missing Values in City Data:\n", city_data.isnull().sum())
print("Missing Values in Customer Data:\n", customer_data.isnull().sum())
print("Missing Values in Transaction Data:\n", transaction_data.isnull().sum())
```

```
# Convert 'Date of Travel' to datetime in Cab_Data
```

```
cab_data['Date of Travel'] = pd.to_datetime(cab_data['Date of Travel'], origin='1899-12-30', unit='D')
```

```
# Standardize city names
cab_data['City'] = cab_data['City'].str.strip()
city_data['City'] = city_data['City'].str.strip()

# Check for duplicates
print("Duplicates in Cab Data:", cab_data.duplicated().sum())
print("Duplicates in City Data:", city_data.duplicated().sum())
print("Duplicates in Customer Data:", customer_data.duplicated().sum())
print("Duplicates in Transaction Data:", transaction_data.duplicated().sum())
```

```
➦ Missing Values in Cab Data:
  Transaction ID    0
  Date of Travel    0
  Company           0
  City              0
  KM Travelled      0
  Price Charged     0
  Cost of Trip      0
  dtype: int64
Missing Values in City Data:
  City              0
  Population        0
  Users             0
  dtype: int64
Missing Values in Customer Data:
  Customer ID       0
  Gender            0
  Age               0
  Income (USD/Month) 0
  dtype: int64
Missing Values in Transaction Data:
  Transaction ID    0
  Customer ID       0
  Payment_Mode      0
  dtype: int64
Duplicates in Cab Data: 0
Duplicates in City Data: 0
Duplicates in Customer Data: 0
Duplicates in Transaction Data: 0
```

```
# Merge Cab_Data with Transaction_ID
merged_data = pd.merge(cab_data, transaction_data, on='Transaction ID')

# Merge with Customer_ID
master_data = pd.merge(merged_data, customer_data, on='Customer ID')

# Merge with City
master_data = pd.merge(master_data, city_data, on='City')

# Display the structure of the final master dataset
print("Master Dataset:")
print(master_data.info())
```

```
➦ Master Dataset:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 359392 entries, 0 to 359391
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Transaction ID        359392 non-null  int64
 1   Date of Travel        359392 non-null  datetime64[ns]
 2   Company               359392 non-null  object
 3   City                  359392 non-null  object
 4   KM Travelled          359392 non-null  float64
 5   Price Charged         359392 non-null  float64
 6   Cost of Trip          359392 non-null  float64
 7   Customer ID           359392 non-null  int64
 8   Payment_Mode          359392 non-null  object
 9   Gender                359392 non-null  object
10   Age                  359392 non-null  int64
11   Income (USD/Month)    359392 non-null  int64
12   Population            359392 non-null  object
13   Users                 359392 non-null  object
dtypes: datetime64[ns](1), float64(3), int64(4), object(6)
memory usage: 38.4+ MB
None
```

✓ Revenue and Margin Analysis by Company

Understanding which company generates more revenue and better margins is critical for investment decisions.

```
# Summary statistics
print(master_data.describe())


# Distribution of revenue by company
master_data.groupby('Company')['Price Charged'].sum().plot(kind='bar', title="Revenue by Company")
```

	Transaction ID	Date of Travel	KM Travelled
count	3.593920e+05	359392	359392.000000
mean	1.022076e+07	2017-08-17 01:37:55.042293760	22.567254
min	1.000001e+07	2016-01-02 00:00:00	1.900000
25%	1.011081e+07	2016-11-23 00:00:00	12.000000
50%	1.022104e+07	2017-09-10 00:00:00	22.440000
75%	1.033094e+07	2018-05-12 00:00:00	32.960000
max	1.044011e+07	2018-12-31 00:00:00	48.000000
std	1.268058e+05	NaN	12.233526

	Price Charged	Cost of Trip	Customer ID	Age
count	359392.000000	359392.000000	359392.000000	359392.000000
mean	423.443311	286.190113	19191.652115	35.336705
min	15.600000	19.000000	1.000000	18.000000
25%	206.437500	151.200000	2705.000000	25.000000
50%	386.360000	282.480000	7459.000000	33.000000
75%	583.660000	413.683200	36078.000000	42.000000
max	2048.030000	691.200000	60000.000000	65.000000
std	274.378911	157.993661	21012.412463	12.594234

	Income (USD/Month)
count	359392.000000
mean	15048.822937
min	2000.000000
25%	8424.000000
50%	14685.000000
75%	21035.000000
max	35000.000000
std	7969.409482


<Axes: title={'center': 'Revenue by Company'}, xlabel='Company'>



```
import matplotlib.pyplot as plt
import seaborn as sns
```

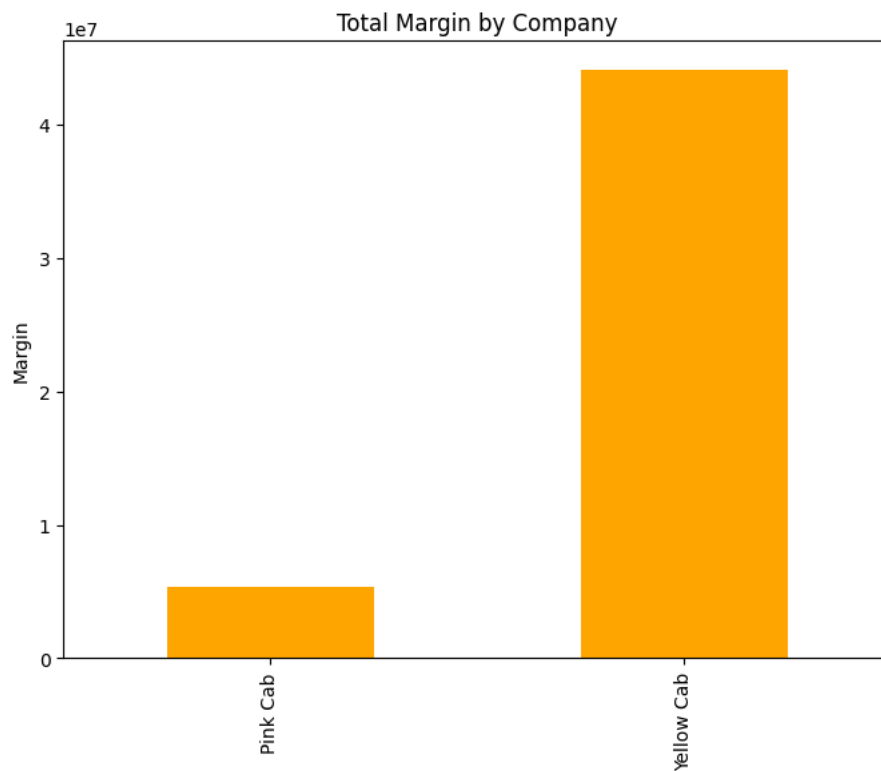
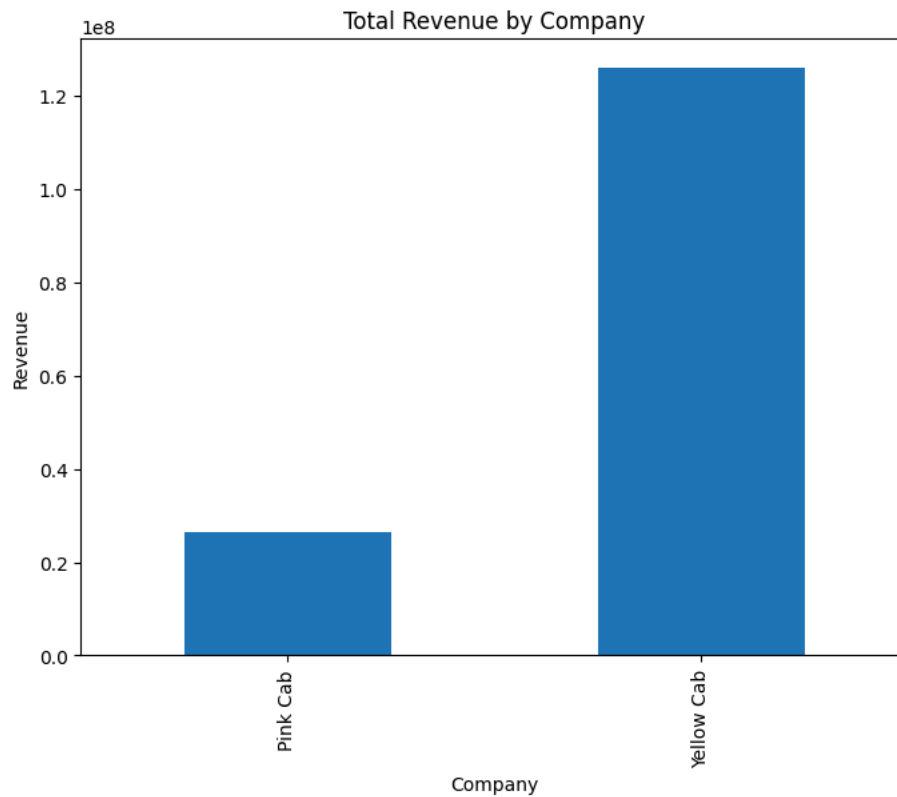
```
# Extract year and month
master_data['Year'] = master_data['Date of Travel'].dt.year
master_data['Month'] = master_data['Date of Travel'].dt.month
```

```
# Monthly revenue trend
monthly_revenue = master_data.groupby(['Company', 'Year', 'Month'])['Price Charged'].sum().unstack(level=0)
monthly_revenue.plot(kind='line', title="Monthly Revenue Trend")
plt.show()
```



```
# Total revenue by company
revenue_by_company = master_data.groupby('Company')['Price Charged'].sum()
revenue_by_company.plot(kind='bar', title="Total Revenue by Company", figsize=(8, 6))
plt.ylabel("Revenue")
plt.show()
```

```
# Total margin by company
master_data['Margin'] = master_data['Price Charged'] - master_data['Cost of Trip']
margin_by_company = master_data.groupby('Company')['Margin'].sum()
margin_by_company.plot(kind='bar', title="Total Margin by Company", color="orange", figsize=(8, 6))
plt.ylabel("Margin")
plt.show()
```




```
# Average income by company
avg_income = master_data.groupby('Company')['Income (USD/Month)'].mean()
print("Average Income by Company:\n", avg_income)
```

```
# Age distribution
sns.histplot(master_data['Age'], bins=10, kde=True)
plt.title("Age Distribution of Customers")
plt.show()
```

```

Average Income by Company:
Company
Pink Cab      15059.047137
Yellow Cab    15045.669817
Name: Income (USD/Month), dtype: float64

```



✓ City-Wise Performance Analysis

Analyze city-wise revenue and the number of users to identify high-performing regions.

✓ Insights from Seasonality in Revenue

1. **Revenue peaks in December:** This suggests a strong seasonal trend, likely influenced by holidays and winter demand.
2. **Lower revenue in early months (January, February):** Possible reduced demand post-holiday season.
3. **Recommendation:**
 - Focus marketing efforts during peak months (November-December).
 - Offer promotions or discounts in low-demand months to attract more customers.

```

# City-wise revenue
city_revenue = master_data.groupby('City')['Price Charged'].sum().sort_values(ascending=False)
print("Top Cities by Revenue:\n", city_revenue.head())

```

```

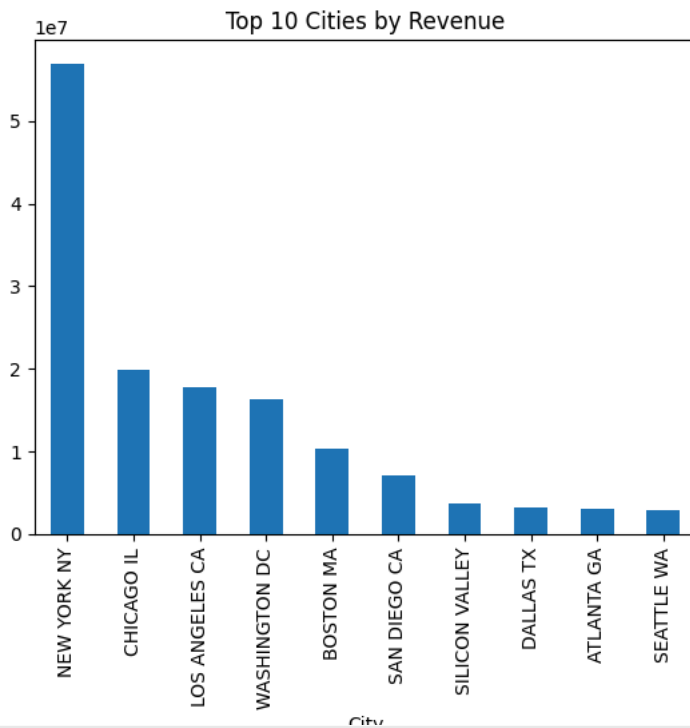
# Visualize top cities by revenue
city_revenue.head(10).plot(kind='bar', title="Top 10 Cities by Revenue")
plt.show()

```

```

Top Cities by Revenue:
City
NEW YORK NY      56954061.67
CHICAGO IL       19841318.52
LOS ANGELES CA   17795624.41
WASHINGTON DC    16366703.83
BOSTON MA        10359755.42
Name: Price Charged, dtype: float64

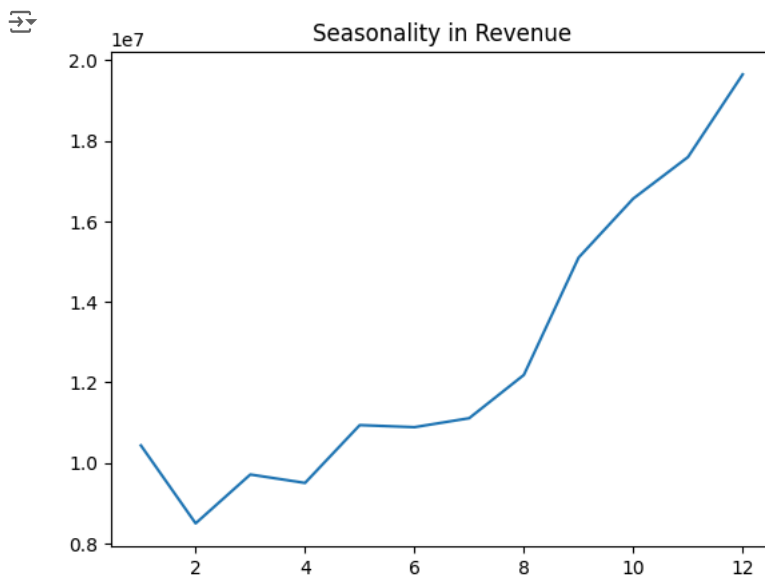
```



```

# Check for seasonality
seasonality = master_data.groupby(master_data['Date of Travel'].dt.month)['Price Charged'].sum()
seasonality.plot(kind='line', title="Seasonality in Revenue")
plt.show()

```




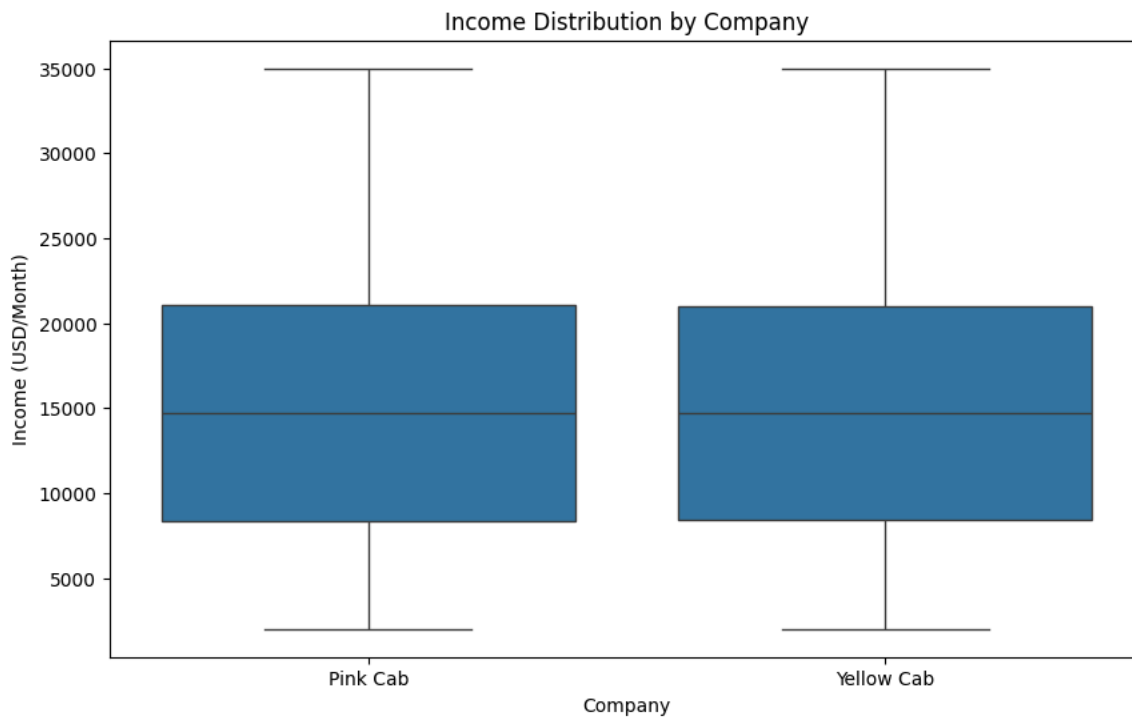
Customer Demographics Analysis

Analyze customer income and age to identify the most valuable customer segments.

```
import seaborn as sns
```

```
# Income distribution by company  
plt.figure(figsize=(10, 6))  
sns.boxplot(x='Company', y='Income (USD/Month)', data=master_data)  
plt.title("Income Distribution by Company")  
plt.show()
```

```
# Age distribution  
plt.figure(figsize=(8, 5))  
sns.histplot(master_data['Age'], bins=15, kde=True, color='purple')  
plt.title("Age Distribution of Customers")  
plt.xlabel("Age")  
plt.show()
```




```
# Check the structure of master_data
print(master_data.info())

# Display the first few rows to confirm the column names
print(master_data.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 359392 entries, 0 to 359391
Data columns (total 17 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Transaction ID         359392 non-null  int64
 1   Date of Travel         359392 non-null  datetime64[ns]
 2   Company                359392 non-null  object
 3   City                  359392 non-null  object
 4   KM Travelled           359392 non-null  float64
 5   Price Charged          359392 non-null  float64
 6   Cost of Trip           359392 non-null  float64
 7   Customer ID            359392 non-null  int64
 8   Payment_Mode           359392 non-null  object
 9   Gender                 359392 non-null  object
10   Age                    359392 non-null  int64
11   Income (USD/Month)     359392 non-null  int64
12   Population             359392 non-null  object
13   Users                  359392 non-null  object
14   Year                   359392 non-null  int32
15   Month                  359392 non-null  int32
16   Margin                 359392 non-null  float64
dtypes: datetime64[ns](1), float64(4), int32(2), int64(4), object(6)
memory usage: 43.9+ MB
None
```

	Transaction ID	Date of Travel	Company	City	KM Travelled	\
0	10000011	2016-01-08	Pink Cab	ATLANTA GA	30.45	
1	10000012	2016-01-06	Pink Cab	ATLANTA GA	28.62	
2	10000013	2016-01-02	Pink Cab	ATLANTA GA	9.04	
3	10000014	2016-01-07	Pink Cab	ATLANTA GA	33.17	
4	10000015	2016-01-03	Pink Cab	ATLANTA GA	8.73	

	Price Charged	Cost of Trip	Customer ID	Payment_Mode	Gender	Age	\
0	370.95	313.635	29290	Card	Male	28	
1	358.52	334.854	27703	Card	Male	27	
2	125.20	97.632	28712	Cash	Male	53	
3	377.40	351.602	28020	Cash	Male	23	
4	114.62	97.776	27182	Card	Male	33	

	Income (USD/Month)	Population	Users	Year	Month	Margin	\
0	10813	814,885	24,701	2016	1	57.315	
1	9237	814,885	24,701	2016	1	23.666	
2	11242	814,885	24,701	2016	1	27.568	
3	23327	814,885	24,701	2016	1	25.798	
4	8536	814,885	24,701	2016	1	16.844	

```
# Inspect the Population and Price Charged columns
print(master_data[['City', 'Population', 'Price Charged']].head())

# Check for non-numeric values in Population
print(master_data['Population'].unique())

# Convert Population to numeric again and handle errors
master_data['Population'] = pd.to_numeric(master_data['Population'], errors='coerce')

# Drop rows where Population or Price Charged is missing
master_data = master_data.dropna(subset=['Population', 'Price Charged'])

# Verify the cleaned data
print(master_data[['City', 'Population', 'Price Charged']].head())
```

```
Empty DataFrame
Columns: [City, Population, Price Charged]
Index: []

Empty DataFrame
Columns: [City, Population, Price Charged]
Index: []
```

```
# Check unique values in Population and Price Charged before cleaning
print("Unique Population Values:")
print(city_data['Population'].unique())
```

```
print("\nUnique Price Charged Values:")
print(cab_data['Price Charged'].unique())
```

```
# Display first rows of city_data and cab_data
print("\nCity Data:")
print(city_data.head())
```

```
print("\nCab Data:")
print(cab_data.head())
```

```
Unique Population Values:
[' 8,405,837 ' ' 1,955,130 ' ' 1,595,037 ' ' 1,339,155 ' ' 1,177,609 '
 ' 1,030,185 ' ' 959,307 ' ' 943,999 ' ' 942,908 ' ' 814,885 ' ' 754,233 '
 ' 698,371 ' ' 671,238 ' ' 631,442 ' ' 629,591 ' ' 545,776 ' ' 542,085 '
 ' 418,859 ' ' 327,225 ' ' 248,968 ']
```

```
Unique Price Charged Values:
[370.95 358.52 125.2 ... 31.49 742.24 620.62]
```

City Data:

	City	Population	Users
0	NEW YORK NY	8,405,837	302,149
1	CHICAGO IL	1,955,130	164,468
2	LOS ANGELES CA	1,595,037	144,132
3	MIAMI FL	1,339,155	17,675
4	SILICON VALLEY	1,177,609	27,247

Cab Data:

	Transaction ID	Date of Travel	Company	City	KM Travelled \
0	10000011	2016-01-08	Pink Cab	ATLANTA GA	30.45
1	10000012	2016-01-06	Pink Cab	ATLANTA GA	28.62
2	10000013	2016-01-02	Pink Cab	ATLANTA GA	9.04
3	10000014	2016-01-07	Pink Cab	ATLANTA GA	33.17
4	10000015	2016-01-03	Pink Cab	ATLANTA GA	8.73

	Price Charged	Cost of Trip
0	370.95	313.635
1	358.52	334.854
2	125.20	97.632
3	377.40	351.602
4	114.62	97.776

```
# Clean 'Population' column in City data
city_data['Population'] = city_data['Population'].str.replace(',', '').astype(float)
```

```
# Clean 'Price Charged' in Cab data
cab_data['Price Charged'] = pd.to_numeric(cab_data['Price Charged'], errors='coerce')
```

```
# Verify cleaned columns
print(city_data.head())
print(cab_data.head())
```

```
City Population Users
0 NEW YORK NY 8405837.0 302,149
1 CHICAGO IL 1955130.0 164,468
2 LOS ANGELES CA 1595037.0 144,132
3 MIAMI FL 1339155.0 17,675
4 SILICON VALLEY 1177609.0 27,247

Transaction ID Date of Travel Company City KM Travelled \
0 10000011 2016-01-08 Pink Cab ATLANTA GA 30.45
1 10000012 2016-01-06 Pink Cab ATLANTA GA 28.62
2 10000013 2016-01-02 Pink Cab ATLANTA GA 9.04
3 10000014 2016-01-07 Pink Cab ATLANTA GA 33.17
4 10000015 2016-01-03 Pink Cab ATLANTA GA 8.73

Price Charged Cost of Trip
0 370.95 313.635
1 358.52 334.854
2 125.20 97.632
3 377.40 351.602
4 114.62 97.776
```

```
# Merge Cab_Data with Transaction_ID
merged_data = pd.merge(cab_data, transaction_data, on='Transaction ID')
```

```
# Merge with Customer_ID
```

```

master_data = pd.merge(merged_data, customer_data, on='Customer ID')

# Merge with City data
master_data = pd.merge(master_data, city_data, on='City', how='left') # Ensure left join for matching

# Group by City for Population and Revenue
city_revenue = master_data.groupby('City')[['Population', 'Price Charged']].sum().reset_index()

# Verify if the data is populated
print(city_revenue.head())

# Plot the data
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.scatterplot(x='Population', y='Price Charged', data=city_revenue)
plt.title("Population vs Revenue")
plt.xlabel("City Population")
plt.ylabel("Total Revenue")
plt.show()

```

```

↔

```

	City	Population	Price Charged
0	ATLANTA GA	6.158086e+09	2980241.72
1	AUSTIN TX	3.419224e+09	1877142.50
2	BOSTON MA	7.392358e+09	10359755.42
3	CHICAGO IL	1.107092e+11	19841318.52
4	DALLAS TX	6.616385e+09	3142429.91

